

# Task2-Wine Quality Prediction

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn import metrics
from sklearn.svm import SVC

from sklearn.linear_model import LogisticRegression

import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: df = pd.read_csv("WineQT.csv")
print(df.head())
```

|       | fixed acidity | volatile acidity | citric acid | residual sugar | chlori |
|-------|---------------|------------------|-------------|----------------|--------|
| des \ |               |                  |             |                |        |
| 0     | 7.4           | 0.70             | 0.00        | 1.9            | 0.     |
| 076   |               |                  |             |                |        |
| 1     | 7.8           | 0.88             | 0.00        | 2.6            | 0.     |
| 098   |               |                  |             |                |        |
| 2     | 7.8           | 0.76             | 0.04        | 2.3            | 0.     |
| 092   |               |                  |             |                |        |
| 3     | 11.2          | 0.28             | 0.56        | 1.9            | 0.     |
| 075   |               |                  |             |                |        |
| 4     | 7.4           | 0.70             | 0.00        | 1.9            | 0.     |
| 076   |               |                  |             |                |        |

|   | free sulfur dioxide | total sulfur dioxide | density | pH   | sulphates |
|---|---------------------|----------------------|---------|------|-----------|
| \ |                     |                      |         |      |           |
| 0 | 11.0                | 34.0                 | 0.9978  | 3.51 | 0.56      |
| 1 | 25.0                | 67.0                 | 0.9968  | 3.20 | 0.68      |
| 2 | 15.0                | 54.0                 | 0.9970  | 3.26 | 0.65      |
| 3 | 17.0                | 60.0                 | 0.9980  | 3.16 | 0.58      |
| . | 11.0                | 34.0                 | 0.9978  | 3.51 | 0.56      |

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1143 entries, 0 to 1142
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1143 non-null   float64
1   volatile acidity       1143 non-null   float64
2   citric acid            1143 non-null   float64
3   residual sugar         1143 non-null   float64
4   chlorides              1143 non-null   float64
5   free sulfur dioxide    1143 non-null   float64
6   total sulfur dioxide   1143 non-null   float64
7   density                1143 non-null   float64
8   pH                    1143 non-null   float64
9   sulphates              1143 non-null   float64
10  alcohol                1143 non-null   float64
11  quality                1143 non-null   int64
12  Id                     1143 non-null   int64
dtypes: float64(11), int64(2)
memory usage: 116.2 KB
```

```
In [5]: df.describe().T
```

Out[5]:

|                             | count  | mean       | std        | min     | 25%       | 50%       | 75%         | max        |
|-----------------------------|--------|------------|------------|---------|-----------|-----------|-------------|------------|
| <b>fixed acidity</b>        | 1143.0 | 8.311111   | 1.747595   | 4.60000 | 7.10000   | 7.90000   | 9.100000    | 15.50000   |
| <b>volatile acidity</b>     | 1143.0 | 0.531339   | 0.179633   | 0.12000 | 0.39250   | 0.52000   | 0.640000    | 1.01000    |
| <b>citric acid</b>          | 1143.0 | 0.268364   | 0.196686   | 0.00000 | 0.09000   | 0.25000   | 0.420000    | 1.06000    |
| <b>residual sugar</b>       | 1143.0 | 2.532152   | 1.355917   | 0.90000 | 1.90000   | 2.20000   | 2.600000    | 15.50000   |
| <b>chlorides</b>            | 1143.0 | 0.086933   | 0.047267   | 0.01200 | 0.07000   | 0.07900   | 0.090000    | 0.19000    |
| <b>free sulfur dioxide</b>  | 1143.0 | 15.615486  | 10.250486  | 1.00000 | 7.00000   | 13.00000  | 21.000000   | 68.00000   |
| <b>total sulfur dioxide</b> | 1143.0 | 45.914698  | 32.782130  | 6.00000 | 21.00000  | 37.00000  | 61.000000   | 289.00000  |
| <b>density</b>              | 1143.0 | 0.996730   | 0.001925   | 0.99007 | 0.99557   | 0.99668   | 0.997845    | 1.00000    |
| <b>pH</b>                   | 1143.0 | 3.311015   | 0.156664   | 2.74000 | 3.20500   | 3.31000   | 3.400000    | 4.01000    |
| <b>sulphates</b>            | 1143.0 | 0.657708   | 0.170399   | 0.33000 | 0.55000   | 0.62000   | 0.730000    | 2.01000    |
| <b>alcohol</b>              | 1143.0 | 10.442111  | 1.082196   | 8.40000 | 9.50000   | 10.20000  | 11.100000   | 14.50000   |
| <b>quality</b>              | 1143.0 | 5.657043   | 0.805824   | 3.00000 | 5.00000   | 6.00000   | 6.000000    | 8.00000    |
| <b>Id</b>                   | 1143.0 | 804.969379 | 463.997116 | 0.00000 | 411.00000 | 794.00000 | 1209.500000 | 1597.00000 |

# EDTA-Analysis

```
In [6]: df.isnull().sum()
```

```
Out[6]: fixed acidity      0
        volatile acidity  0
        citric acid       0
        residual sugar    0
        chlorides         0
        free sulfur dioxide 0
        total sulfur dioxide 0
        density           0
        pH                0
        sulphates         0
        alcohol           0
        quality           0
        Id                0
        dtype: int64
```

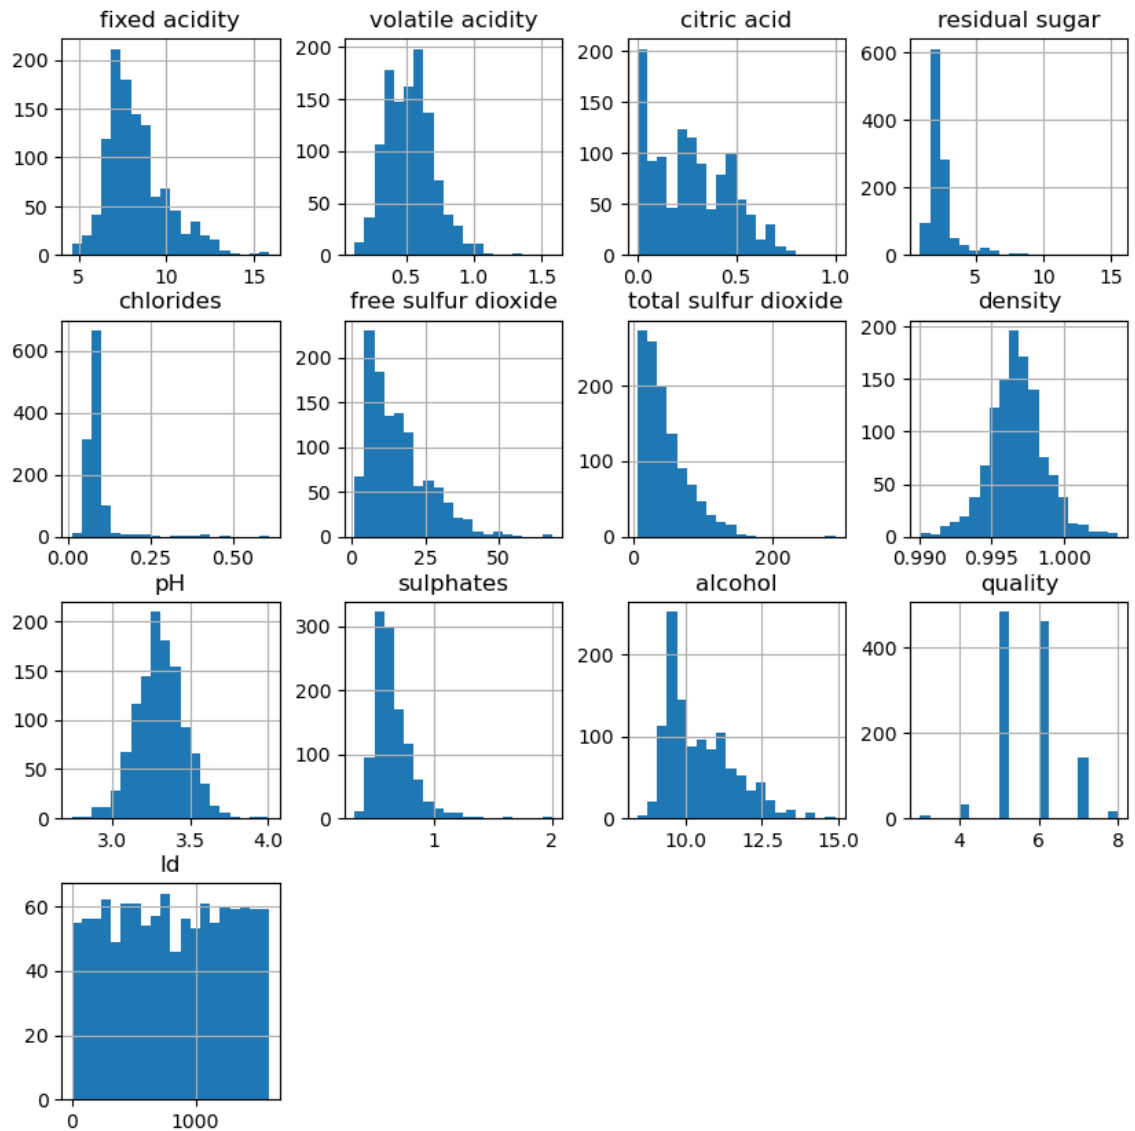
```
In [7]: for col in df.columns:
        if df[col].isnull().sum() > 0:
            df[col] = df[col].fillna(df[col].mean())

df.isnull().sum().sum()
```


```
Out[7]: 0
```

## Histogram


```
In [8]: df.hist(bins=20, figsize=(10, 10))  
plt.show()
```



```
In [9]: plt.bar(df['quality'], df['alcohol'])  
plt.xlabel('quality')  
plt.ylabel('alcohol')  
plt.show()
```



```
In [10]: plt.figure(figsize=(12, 12))
sb.heatmap(df.corr() > 0.7, annot=True, cbar=False)
plt.show()
```



```
In [11]: df = df.drop('total sulfur dioxide', axis=1)
```

```
In [12]: df['best quality'] = [1 if x > 5 else 0 for x in df.quality]
```

```
In [13]: df.replace({'white': 1, 'red': 0}, inplace=True)
```

```
In [14]: features = df.drop(['quality', 'best quality'], axis=1)
target = df['best quality']

xtrain, xtest, ytrain, ytest = train_test_split(
    features, target, test_size=0.2, random_state=40)

xtrain.shape, xtest.shape
```

```
Out[14]: ((914, 11), (229, 11))
```

```
In [15]: #Normalizing the data
norm = MinMaxScaler()
xtrain = norm.fit_transform(xtrain)
xtest = norm.transform(xtest)
```

```
In [17]: models = [LogisticRegression(),SVC(kernel='rbf')]

for i in range(3):
    models[i].fit(xtrain, ytrain)

    print(f'{models[i]} : ')
    print('Training Accuracy : ', metrics.roc_auc_score(ytrain, models[i].p
    print('Validation Accuracy : ', metrics.roc_auc_score(
        ytest, models[i].predict(xtest)))
    print()
```


```
LogisticRegression() :
Training Accuracy : 0.7546950559364851
Validation Accuracy : 0.7255154639175256
```

```
SVC() :
Training Accuracy : 0.7648213641284736
Validation Accuracy : 0.7358247422680412
```

```
-----
-
IndexError                                Traceback (most recent call las
t)
~\AppData\Local\Temp\ipykernel_23416\723548108.py in <module>
      2
      3 for i in range(3):
----> 4     models[i].fit(xtrain, ytrain)
      5
      6     print(f'{models[i]} : ')
```

**IndexError:** list index out of range

```
In [18]: metrics.plot_confusion_matrix(models[1], xtest, ytest)
plt.show()
```



```
In [19]: print(metrics.classification_report(ytest,
                                              models[1].predict(xtest)))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.68      | 0.72   | 0.70     | 97      |
| 1            | 0.79      | 0.75   | 0.77     | 132     |
| accuracy     |           |        | 0.74     | 229     |
| macro avg    | 0.73      | 0.74   | 0.73     | 229     |
| weighted avg | 0.74      | 0.74   | 0.74     | 229     |

```
In [ ]:
```